

Apportioning Responsibilities

Author(s): Robert E. Goodin

Source: *Law and Philosophy*, Aug., 1987, Vol. 6, No. 2 (Aug., 1987), pp. 167-185

Published by: Springer

Stable URL: <https://www.jstor.org/stable/3504573>

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <https://about.jstor.org/terms>



JSTOR

Springer is collaborating with JSTOR to digitize, preserve and extend access to *Law and Philosophy*

APPORTIONING RESPONSIBILITIES

Occasionally, one individual is said to bear exclusive responsibility for some particular state of affairs. More commonly, several individuals are said to share responsibility, and indeed to share it differentially. Responsibility, in the typical case, is neither exclusive nor equal. There are ordinarily some people who are thought to bear rather more responsibility, and others who are thought to bear rather less, for one and the same state of affairs.¹

That is simply to say that we characteristically find ourselves *apportioning* responsibility between different agents. In itself, that observation seems unexceptional and unexceptionable. Upon further inspection, however, it turns out that that observation is not nearly so idle as it seems. That we apportion responsibility, and how we do so, in the end tells us something important about the nature of responsibility itself. For it is only on certain models of responsibility – the standard model, extraordinary as it seems, not being among them – that that practice makes sense.

The standard model of responsibility (championed variously by lawyers, churchmen and Kantian philosophers) is essentially a mechanism for fixing credit or, more commonly, blame for certain sorts of states of affairs. Hence I shall dub this model one of ‘blame-responsibility’. People are paradigmatically held responsible for flaws in their characters and for the outcomes (especially the harms) that

¹ Neither, where responsibility is shared among several people, do their fractional responsibilities necessarily sum to exactly one. We can (and sometimes should) follow the principle of tort law, whereby people are held jointly and severally liable for the results of their concerted action; then each is fully responsible for the entire amount of damages caused by all. On this, see Robert E. Goodin, *Protecting the Vulnerable* (Chicago: University of Chicago Press, 1985) p. 135 and references therein.

flow from them. People who harbour evil intentions, producing harms on purpose and by design, are charged with the gravest form of responsibility, marked off by the criminal law and its moral equivalents; those who are merely negligent, unintentionally but recklessly causing harm, are charged with lighter forms of responsibility, marked off by the civil law and its moral analogues. On both levels, the central concern of advocates of the model of blame-responsibility is with 'what sort of person one is', as revealed through one's actions and what (motives, thoughtlessness, etc.) can be seen lying behind them.

Presumably ascribing responsibility is always essentially a matter of pointing fingers. What distinguishes the standard model of responsibility from the one that I shall here be advocating is the *purpose* for which fingers are being pointed. With the standard model of responsibility, fingers are being pointed for purposes of fixing blame. With my model of responsibility, fingers are being pointed for purposes of assigning duties, jobs or (generically) tasks. Hence I shall dub my model one of 'task-responsibility'. In it, questions of 'who has been assigned what tasks?' become truly central to the business of ascribing responsibilities. The difference between 'my job' and 'your job', on this account of responsibilities, looms particularly large.

Of course, it would make no sense to assign people task-responsibilities if they were not also going to be held to account for how well or badly they performed those tasks and discharged those responsibilities. Therefore, task-responsibilities inevitably entail correlative blame-responsibilities. But the blame-responsibilities entailed by task-responsibilities are different, both in substance and in form, from those that are implied by the standard model of blame-responsibilities. The former put questions of consequences to the fore, the latter questions of character to the fore.

It is the argument of this paper that these differences matter. In it, I shall be trying to show that my model of task-responsibility makes better sense of the way we apportion differential responsibilities than does the standard, character-based model of blame-responsibility.

I.

Consider this distressingly realistic pair of scenarios. In the first, a terrorist gang has exploded a bomb hidden in some rubbish at the back of a busy store, killing several innocent shoppers and injuring several more. How should we apportion responsibilities for this damage? Intuitively, and after considerable reflection as well, we would ordinarily say something like this. The terrorist who master-minded the attack bears greatest responsibility for its consequences. The terrorist who actually detonated the bomb bears a greater responsibility than the member of the gang who merely planted the bomb in the store and did nothing more. All of those agents bear a greater responsibility than the store detective whose perfunctory search of the premises that morning failed to uncover the bomb, but his responsibility is greater than that of the janitor who failed to clear the rubbish the night before. All of them, however, bear greater responsibility than the politicians and the electorate whose decisions gave the terrorist gang a justifiable grievance.

Or, for a second scenario, imagine a particularly messy divorce that has so alienated a mother from the child left in her custody that she neglects it and it dies in consequence. How should we apportion responsibilities for this death? Intuitively, and after considerable reflection as well, we would ordinarily say something like this. The brunt of responsibility for the child's death must fall upon the neglectful mother. But a substantial, though lesser, measure of responsibility falls upon the ex-husband who, had he exercised court-awarded visiting rights, could have prevented the tragedy. A lesser, though nonetheless substantial, measure of responsibility falls to the Public Health Visitor, who should have been calling at the family's house to check on the health of infants in general. And some modicum of responsibility falls to all those (the child's neighbours, grandparents, etc.) who might reasonably have been expected to observe the child's distress and report it to the proper authorities.

What is revealed by these cases, and hundreds of others like them, is that we ordinarily want to make fairly fine-grained judgments about the relative degrees of responsibility borne by various actors in

the sequence of events that conspired to produce any particular state of affairs. But what sense can be made of these judgments? On what basis can we reasonably apportion responsibility?

II.

Advocates of the model of blame-responsibility offer a bifurcated response to such questions. First and foremost, anyone possessed of *mens rea* – of evil intent – is criminally guilty, and bears proportionately greater responsibility than anyone in whom this element of intentionality is missing. Second, and very much secondarily, anyone guilty of negligence bears some responsibility; just how much depends upon the magnitude of the harm and the probability that the negligent agent should have foreseen its occurring, given his actions.

At first blush, this account might seem broadly right. Surely it is correct to say that, in the case of the bomb blast, all the terrorists (who harbour evil intent) come in for greater shares of responsibility than do watchmen, janitors or politicians (who are guilty of mere negligence). And it is probably also right to say that everyone whom we want to hold responsible to any extent at all, in that case or any other, is guilty (at the very least) of negligence of some sort of another.

Upon closer analysis, however, this initial plausibility fades. Questions arise concerning the adequacy of this model on three levels: its account of first-class, intentions-based responsibilities; of second-class, negligence-based responsibilities; and of the relationship between these two classes of responsibilities. On none of these points is the model of character-based blame-responsibility wholly convincing.

A. At the higher level – that of intentions-based responsibilities – the main problem is that this model of blame-responsibility does not allow us to make any very fine distinctions between different people's differential responsibilities. Dealing strictly in terms of intentions, we can make only the grossest distinction between those who were party to the terrorists' conspiracy and those who were not. Suppose the terrorist gang held long and intense discussions about its plan of action

and the anticipated consequences. Suppose every member of the gang agreed to the plan, not just in the knowledge that innocent shoppers would die, but actually welcoming that result as part of some larger political programme. Then in looking at intentions alone, we would be unable to distinguish between the responsibility borne by the terrorist who actually detonated the bomb and that borne by the one who merely planted it.

Now, perhaps that is the right answer after all. Maybe we were simply wrong in attempting to impose any distinction here in the first place. But there is something undeniably disconcerting about a theory that obliges us to say that there is no moral difference whatsoever between central actors (those who masterminded the plot or pushed the detonator) and bit players (those merely buying explosives or acting as decoys while they were planted), just because they were all co-conspirators sharing the same intention. And there is something odder still in saying that the same responsibility – *exactly* the same – is borne by someone who, although participating wholeheartedly (but not figuring especially centrally) in the planning sessions and thus sharing the group's intentions, fell ill and could take no part whatsoever in carrying out the plot. A theory that would hold equally responsible those who played virtually no role in the murders and those who were most instrumental in it, just on account of some shared intentions, is a very queer doctrine indeed.

What has caused all this trouble is something deep within the notion of 'intention'. Intentionality itself does not admit of degrees. On the most plausible analysis, some particular outcome is said to be intended by you if and only if that outcome is necessarily part of what it is to accomplish your desired ends.² Thus, the death of many innocent civilians must be regarded as an intended (because intrinsic) part of the American plan in dropping the atomic bomb on Hiroshima, rather than as an incidental by-product of American

² This discussion is based on Gerald Dworkin, 'Nuclear Intentions', *Ethics*, 95 (April 1985): 445–60, 453, who is borrowing in turn from Judith Jarvis Thomson, *Acts and Other Events* (Ithaca, N.Y.: Cornell University Press, 1977), chapter 19.

intentions to bring the war to a speedy conclusion. And this can best be seen by noticing that if all noncombatants had been evacuated from Hiroshima before the bombing, then the Americans would have had to choose as their target another city that was still functioning as a 'normal' city. (Otherwise, the bombing would amount to the sort of mere 'demonstration' that American planners explicitly rejected.) An effect is properly regarded as a foreseen but unintended by-product of our intentional action only if, were the effect in question somehow removed, we would still want to persist in the same plan of action. If that is not the case, then that effect must be regarded as an intrinsic part of our intentions.

Understood that way, an outcome is either intended or not. The question is not *how much* you would have to change your plans if that consequence were to be avoided; the question is merely *whether* your plans would have to be changed. On its face, then, this analysis seems to offer no way to accommodate 'degrees' into talk of intentionality. Yet without that, it seems impossible to accommodate our desire to talk about 'degrees of responsibility'. We would apparently be unable to 'apportion' responsibility in any meaningful sense: we would just have to assign whole masses of people to one category or another. It would be impossible to say that one person is 'more responsible', or 'bears a greater responsibility', for an outcome than a variety of other people who may share the same intention but who performed a variety of really quite different acts.³

Several manoeuvres are available to advocates of intentions-based

³ Similarly, if an agent's 'intent' is defined in terms of the 'descriptions' under which the agent has chosen the action, then any particular aspect of the action must be either intended (i.e., figure in the agent's own act-description, under which the agent chose the act) or else not intended (i.e., not so figure). Talk of 'degrees of intentionality' must, in these terms, mean something like: 'This act-description is relatively nearer to (or further from) the agent's own'. But, in these terms, there is no room for such a metric. Either the act-description is the agent's own (in which case the act was fully intended under that description), or else it is not (in which case it was not intended at all under that description). This seems to be the upshot of the views of, e.g., Donald Davidson, *Essays on Actions and Events* (Oxford: Clarendon Press, 1980).

accounts of responsibility in trying to evade these criticisms. One of them is to point to a variety of conditions that diminish (or perhaps extinguish altogether) one's responsibility for an action. These work by qualifying (or, in the limiting case, rebutting altogether) the claim that the action was indeed performed intentionally. In terminology once again borrowed from the law, these qualifications might range from 'pleas in mitigation' to fully-fledged 'excuses for actions'.

It may well be true that we want to accept these sorts of qualifications whenever such cases arise. The question here, however, is simply whether qualifications such as these can fully account for the systematic variation in the responsibilities of various members of our hypothetical terrorist gang. The answer to that question seems certainly to be 'no'. While in any particular instance it may be true that the bomb-planter has better excuses than the bomb-detonator, there is no reason to believe that that will always be the case. Indeed, sometimes the opposite is bound to be the case – sometimes it will be the bomb-detonator who is participating under duress, to save his children who are being held hostage by the terrorist gang. Talk of excusing, mitigating or extenuating circumstances seems constitutionally incapable of accounting for the *systematically* greater responsibility of the one party. By its very nature, all such talk must deal in terms of particular features of particular circumstances, and these are bound to vary.

A second evasive manoeuvre is available to advocates of intentions-based accounts of responsibility. This, to continue still further the legal analogy, is to suggest that those with lesser responsibility have in effect 'copped a plea'. That is to say, they plead guilty to a lesser charge, in exchange for being let off the more serious one. The strategy here is to subdivide the plan and assign responsibility for different parts of its formulation and execution to different agents. Thus, the bomb-planter might be thought guilty merely of 'conspiracy to murder' or of 'endangering life', whereas the bomb-detonator might be thought guilty of murder in the first degree.⁴

⁴ Something like this is suggested in Joel Feinberg's handling of the problem in *Doing and Deserving* (Princeton: Princeton University Press, 1979) pp. 244–6.

One flaw in that approach to the problem of differential blame-responsibility is this. If we stick exclusively to the story about intentions, and if intentions are conceived in the terms described above, then you would have to plead *fully* guilty to the lesser charge and be released *entirely* from responsibility on the graver one. That is to say, any gradations of responsibility must be located in gradations of more or less serious offences, not in differential degrees of responsibility for the same offence. Whether or not we can in this way get all the shading we want and make all the fine-grained distinctions we think we should in our system of responsibility thus depends upon just how many different offences we can distinguish.

Perhaps it might be logically possible to invent enough distinct offences to do the job. But it seems empirically unlikely that there are enough already on the statute books or the equivalent moral registers. To return to the previous example, we think different members of the terrorist gang bear differential responsibilities for the murders, even though we cannot find different moral or legal pigeonholes for them. The decoy or purchaser of explosives bears greater responsibility than the financial backers or passive members of the gang, for example, but it is hard to think how to describe all those differential responsibilities in terms of responsibilities for different moral or legal offences. Hence it seems that we are holding them differentially responsible for the same thing, rather than holding them responsible for different things.⁵

A third and final evasive manoeuvre available to advocates of intentions-based models of blame-responsibility is to reformulate the notion of 'intentions'. On this revised account, what each participant in the terrorist conspiracy intends – but all that any of them intends – is to play one's *own* part in the plan. Crucially, on this account, none

⁵ This is not to deny that lawyers and moralists alike hold people differentially responsible for identical offences. Clearly they do, as is powerfully reflected in, e.g., sentencing decisions. My point is merely that when judges and moralists make such distinctions, they do so only by abandoning the single-minded focus upon intentions that officially characterizes this half of the model of blame-responsibility.

of them (with the possible exception of the mastermind himself) would truly intend the outcome of their combined actions. Differential responsibilities would then follow straightforwardly. Those who intend lesser contributions are less responsible for the overall outcome than those who intend greater contributions.

This might be a reasonable account of responsibilities for *accidents*, where the intended actions of various agents interact in ways that none of them intend in order to produce the outcome that all regard as unfortunate. It is not, however, a plausible account of conspiratorial intentions. There, the interactions between their intentional actions are deliberate, not accidental or incidental. The intentions of each conspirator ‘track’ the larger goal of the conspiracy as a whole.⁶ If any one of the conspirators alters in any way his intention to participate in the conspiracy, all the other conspirators would (if they know of his altered intentions) consider what alterations are therefore required in their own intentions in order to achieve the aims of the conspiracy. In light of this, no true co-conspirator could plausibly claim to be performing intentionally just his own component of the overall plan. Rather, each intends the whole.

B. The second half of the model of blame-responsibility assigns responsibilities of a weaker sort to people on the basis of their negligence. The problem with this half of the model is not that it lacks an account of differential responsibilities, as does the other intentions-based half. Clearly, different people do end up with different degrees of negligence-based responsibility. The question is whether the right people end up being assigned the right amounts.

The extent of negligence-based blame-responsibility varies explicitly with (a) the magnitude of the harm that occurs and (b) the foreseeable probability of its occurring, given the negligent party’s actions or omissions. Thus, it is easy to see how one person can be said to be more negligent (and hence more responsible, in that sense) than an-

⁶ Robert Nozick, *Philosophical Explanations* (Cambridge: Harvard University Press, 1981) pp. 317–26.

other. On the analysis here in view, that amounts merely to saying that either the magnitude or the probability or both was greater in his case than in the other's case.

Sometimes that will lead to results that are broadly consistent with our intuitive allocation of responsibilities, but often it will not. In the case of the terrorist attack, we are inclined to hold the security guard whose search failed to turn up the bomb more responsible for the ensuing mayhem than we are the janitor who merely failed to collect on time the rubbish in which the bomb was hidden, and that may be because security guards hired to search for bombs are more likely to come across them than are janitors. Even there, however, there is some room for doubt. While that is possibly true as a general rule about guards and janitors, it may well be untrue for janitors working in stores that also have to hire security guards, for example.

Further down the ladder of responsibility, however, doubts about the model's adequacy really begin to mushroom. Consider the responsibilities of politicians and electorates who vote for policies that are sure to provoke terrorist bombing campaigns. We are inclined to say, both at first blush and upon reflection as well, that they bear less responsibility for the deaths than the guard or the janitor. But their actions indisputably put more people more at risk of death from terrorist bombs than the actions of either of those store employees. Even though both the magnitude of harm and the probability of harm increase more as a result of their actions, however, we want to hold them less responsible. The negligence model, analyzing blame-responsibility purely as an amalgam of those two factors, is at a loss to explain why.

Similarly in the case of child-neglect, we are inclined to hold the Health Visitor more responsible for the child's death, even though that official seems less 'negligent' (on the magnitude/probability calculus of negligence) than the child's neighbours. The neighbours, after all, could observe at close quarters a consistent pattern of neglect over a fairly protracted period. Surely there was a higher probability of a more disastrous outcome resulting from their failure to report the case to the proper authorities than from the Health

Visitor's failure to perform the sort of cursory inspection that is all that the law requires. Yet we hold the Health Visitor more responsible.

The upshot seems to be that, while the model of negligence-based blame-responsibility is indeed scalar, it operates on the wrong scale. How negligent one has been, and how much blame-responsibility one will therefore have to bear for the outcome, is analyzed within that model in terms of how 'irresponsibly' one has behaved. Irresponsible behaviour is analyzed in terms of 'thoughtlessness', with how much thought one should have given some action being assessed in terms of the stakes involved, their magnitude and the foreseeable probability that the loss will indeed be incurred. The examples just offered suggest that there is something more to the story than magnitudes and foreseeable probabilities of harm, however. Even in cases of pure negligence (where there is no hint of evil intent), responsibilities are in fact apportioned in ways other than the probability/magnitude model would suggest. It is not just reflection upon people's characters – not just a disdain for 'thoughtlessness' (of that sort, anyway) – that guides our judgments about their responsibilities.

C. Finally, there are problems with the account offered by the blame-responsibility model of the relationship between these two levels of responsibility, one intentions-based and the other negligence-based. Both Kantian philosophy and the legal analogy seem to suggest that intentions-based responsibilities are an order of magnitude more serious than negligence-based responsibilities. The relationship between the two, on the blame-responsibility model, seems to be virtually lexicographical.

Broadly speaking, it is probably right to say that we hold people more responsible for harms they have done intentionally than for harms they have done accidentally. But anything like a lexicographical ranking badly overstates the extent of the priority which intentions enjoy in our thinking about responsibilities. At the margins, we can easily conceive of cases of really *gross* negligence that we would con-

sider more blameworthy than intentional minor contributions to a harmful outcome.⁷

For an example, reflect again upon the case of the terrorist bombing. Conjure up an image of a security guard who did not merely overlook the bomb while doing his ordinary morning rounds, but was instead appallingly negligent. Suppose this was an important anniversary that the terrorists regularly celebrated with bomb attacks; suppose this particular store was a favourite target, having been bombed often before; suppose the previous bombs had all been planted in just the same sort of rubbish piles in just such a corridor. Suppose the security guard in question knew all this. Yet knowing this, and knowing that he will be the only guard to partol the corridor, he gives the rubbish in which the bomb is hidden only a cursory kick and reports back that there is no bomb there. Finally, suppose that it is a case of pure, unmitigated negligence: the guard is neither a secret sympathizer with the terrorists, nor is he prevented from doing his duty properly by paralyzing fear or threats against the lives of his children or such like. Although the guard is merely negligent in the performance of his duty, he is so *grossly* negligent that the responsibility he bears for the shoppers' deaths is at least as great as that borne by some other members of the terrorist gang who made only very minor contributions (e.g., acting as decoy while the bomb was planted) to the success of the bomb attack.

One such case is enough to show that intentions-based responsibilities do not enjoy lexicographical priority over negligence-based ones. There may be some scope for argument over whether this reversal of the 'natural' order is merely a marginal phenomenon. I am inclined to believe that the margins within which it occurs are really rather wide. But there is no need to enter into that argument here. Advocates of models of blame-responsibility, if they admit that this occurs even at the margins, are then obliged to produce a scalar for intentions-based responsibilities, and to produce some mechanism for mapping this

⁷ Indeed, under English law '*gross negligence*' can sometimes give rise to criminal liability; H. L. A. Hart, *Punishment and Responsibility* (Oxford: Clarendon Press, 1968) p. 137.

scale onto the scalar for negligence-based responsibilities to tell us how to trade the one off for the other. Only a tight lexicographical ranking of intentions-based over negligence-based responsibilities would have allowed them to dispense with either (much less both) of these elements in their model. But any such tight lexicographical ranking fails to do justice to our ordinary judgments (intuitive and considered ones, alike) about how to apportion responsibilities in some important cases.

III.

All those are problems that arise within the model of blame-responsibility. There, assigning responsibilities is seen as being essentially a matter of assigning blame. On the model of task-responsibility, in contrast, assigning responsibility amounts essentially to assigning duties and jobs.⁸ Different people have different responsibilities, *ex ante*, because they are allocated different duties and tasks. And people bear differential *ex post* responsibilities for outcomes, on this account, depending on the role that they played or should have played, pursuant to those *ex ante* task-responsibilities, in producing or averting those outcomes.⁹

There are, broadly speaking, two ways in which people might be said to bear different task-responsibilities. In the first, they will have

⁸ Strictly speaking, 'task' is a generic term covering both cases. Colloquially, however, it is sometimes more comfortable to speak of as 'duties' what might more accurately be termed 'negative task responsibilities'. These are responsibilities for making sure that certain things *not* occur. Thus, it is the duty (negative task responsibility) of everyone to make sure they do not cause others harm. (This is what I have elsewhere called a 'self-supervisory responsibility'; see Goodin, 'Responsibilities', *Philosophical Quarterly* 36 (1986): 50–56.) It is the duty (negative task responsibility) of the guard to make sure the store is not bombed. And so on.

⁹ Failure to discharge a task-responsibility might consist in either an act or an omission, depending upon what the task-responsibility required of one. But even those who attach great moral importance to the distinction between acts and omissions in general will standardly concede that there is not necessarily any moral difference between them where the agent is under some *special*

been assigned responsibility for performing different duties or tasks. In the second, they will have been assigned different responsibilities *vis-a-vis* the same duty or task. An important example of the latter is the way in which we assign some people (e.g., godparents) secondary, back-up responsibilities for stepping in and doing certain things (e.g., raising their godchildren) when those with primary responsibility (e.g., their natural parents) fail to discharge it.

I shall say little about the bases upon which these task-responsibilities get assigned to the particular people they do.¹⁰ No doubt part of the story – no doubt a large part of it – has to do with people's differential capacities for performing the tasks and duties at issue. Perhaps some people are peculiarly well-situated to act, or to know if and when action might be necessary. Or perhaps they are particularly blessed with the peculiar sorts of resources (be they material or affective) which the task requires. Or perhaps, in cases where it is crucial to the success of the enterprise that some people act and that others stay well out of their way, they possess some peculiar attribute that makes it 'obvious' to themselves and to others that they are the ones who should act. All of these are essentially consequentialistic considerations, assigning task-responsibilities to people on the grounds that they are particularly able to discharge them. But there may also be other non-consequentialistic considerations at work in the allocation of task-responsibilities. Whether or not there are is a question that, for present purposes, can safely be left open.

My main concern here is instead to show that the model of task-responsibility can offer a plausible account of the way in which we do, in practice, apportion responsibilities among different people. Specifically, the model of task-responsibility can successfully avoid the pitfalls found, in Section II above, with the model of blame-responsibility.

duty or responsibility. Thus, a sleeping sentry is as blameworthy for any harm that results from his negatively omitting to do his duty as is a drunk driver for any harm that results from his actively violating his. On this, see Goodin, *Protecting the Vulnerable*, chapter 2 and references therein.

¹⁰ Cf. Goodin, *Protecting the Vulnerable*, especially chapter 5.

The first problem, discussed in Section II.A, is that the blame-responsibility model, because of its narrow focus upon people's intentions in allocating top-level responsibilities, cannot account for the differential responsibilities borne by people sharing identical intentions but taking very different actions. On the model of task-responsibility, there is no such difficulty in accounting for the differential responsibilities of the various members of the terrorist gang (the mastermind, the detonator, the decoy, the passive member, etc.). The task-responsibility analysis would start from the proposition that each of these people had, *ex ante*, a duty (i.e., a negative task responsibility) to make sure that they did not cause harm to others. From there, it would go on to observe that all these various agents have failed, more or less badly, to discharge that task. How 'badly' they have failed depends, in turn, on just how large a causal contribution each has made to the proscribed outcome.¹¹ The relative strength of each party's task-responsibility for the terrorist bombing thus depends upon the relative importance of the role each played in the overall drama. Lesser degrees of responsibility fall to those (such as passive gang members and decoys) who played what were only very minor roles, in the sense that the basic plot – the overall outcome – would have been little changed if they had been omitted altogether.¹²

¹¹ The intentions-based model of blame-responsibility, having to blame someone either fully or not at all for any particular offence, can only attempt to distinguish various different, more-or-less serious offences in order to build in notions of differential responsibility for the terrorist bombing. The model of task-responsibility, in contrast, admits of degrees in its assessment of any particular offence. The duty that has been violated is the same for all the terrorists – viz., the duty to make sure that they do not harm others. Terrorists vary not in the duty that each has violated, but rather in how badly each has violated it.

¹² When, as occasionally happens, the success of bombing campaigns depends crucially upon the work of a decoy or the moral support of otherwise passive supporters, then both ordinary intuitions and the model of task-responsibility would indeed lay a much larger share of responsibility upon these participants. If all were equally necessary for the success of the mission, then all would be equally responsible for it.

Greater degrees of responsibility fall to those (such as the mastermind and bomb-detonator) who played major, crucial roles in the drama.¹³

Second, the model of blame-responsibility runs into difficulties in analysing lower-level responsibilities purely in terms of negligence, understood as probable harm. To revert to the example of Section II.B, we are often inclined to hold Health Visitors more responsible for failing to report cases of child neglect than we are the child's neighbours, even if in any particular instance the neighbour's inaction is more probable to cause more harm. Why the Health Visitor should be held more responsible in this way is a mystery on the model of blame-responsibility's account of negligence. But it makes perfectly good sense on the model of task responsibility. On that account, the essence of negligence is not probable harm, but rather the shirking of one's duties. It is the job of the Health Visitor, in a way it is not the job of the neighbours, to detect and report cases of child neglect. If the Health Visitor fails in that duty, the Health Visitor therefore bears more responsibility for the subsequent death.

Third, the model of blame-responsibility suffers from problems, discussed in Section II.C, arising from the sharp disjunction in its treatment of higher (intentions-based) and lower (negligence-based) levels of responsibility. In the model of task-responsibility, everything depends upon how important a task you have been assigned and upon how well or badly you perform it. Thus, it is perfectly capable of saying – in a way that the model of blame-responsibility cannot, or can only after considerable twisting and turning – that the grossly negligent guard is more responsible for the innocent shoppers' deaths than is the financier or the perfectly passive member of the terrorist gang. There is no guarantee that there will always be a perfectly continuous measure of role-importance in every instance, of course. But at least there is no built-in guarantee of the sharp and systematic discontinuity that characterises the model of blame-responsibility.

¹³ Assuming the success of the bombing attack will probably depend most crucially upon the successful coordination of the activities of several co-conspirators, the mastermind who performs that coordination role bears the greatest share of responsibility.

Different people have different responsibilities, on the model of task-responsibility, most fundamentally because different people have been assigned different tasks. Blame for failing to discharge those task-responsibilities may be greater or lesser, depending upon how large a role they require one to play in producing the outcome in question.¹⁴ Those with minor roles bear less responsibility than those with major roles.

Blame for failing to discharge one's task-responsibilities may also be greater or lesser depending upon how clear one's task-responsibilities were at the time one was expected to discharge them. Sometimes one is justifiably uncertain whether or not something is one's own responsibility – one's task, duty or job – at all. Other times one is justifiably uncertain as to what particular actions are required in order to produce some particular outcome that one is responsible for making sure is produced. If the *ex ante* perspective of the task-responsibility model is the ruling one, then one's responsibility must be less the less clear it was at the time what one's task-responsibilities were or what they required one to do. And surely this is the right way to think about responsibility. A security guard who failed to search the corridor in which the terrorist bomb was hidden would surely have more success in defending his negligence by saying that he justifiably thought that someone else was supposed to be searching that corridor (the classic excuse for negligence on the task-responsibility model) than he would by saying that he justifiably thought it highly improbable that any very large bomb could have been hidden there (the classic excuse for negligence on the blame-responsibility model).

IV.

There is one final piece of evidence in favour of the task versus blame model of responsibility. It is just this. We can say quite con-

¹⁴ Or, in the case of negative task responsibilities (e.g., to make sure you do not harm others), depending on how large a role one played in producing an outcome which one was duty-bound to avoid producing.

fidently that the terrorist setting off the bomb is more responsible than the terrorist planting it for the deaths of the innocent shoppers, and that both are more responsible than the security guard or the janitor. By the same token, we can say quite confidently that the neglectful mother is more responsible than the slack Health Visitor for the starvation of the infant. But we are hard-pressed to say, in any very precise way at all, exactly *how much* more responsible any of these agents is than the others.¹⁵

Now, there is no reason on either model to suppose that these ratios should be particularly easy to calculate. So the fact that we find the calculations difficult is not necessarily evidence, in and of itself, in support of either model.¹⁶

What does count as decisive evidence in favour of the task-responsibility model is not that we find precise calculations of relative responsibilities difficult, but rather that we find them *dispensable* for such a broad range of practical affairs. For the *ex ante* purposes guiding my choice of actions, I merely need to know whether it is more my responsibility than yours to see to it that some particular task gets done. How much more my responsibility than yours it is, from this *ex ante*, action-guiding perspective, does not matter. Just so long as it is clearly more my task than yours, the allocation of task-responsibilities will be clear enough. The fact that we ordinarily seem so satisfied with ordinal rankings of responsibilities, and so indifferent

¹⁵ That is not to say that we cannot make such judgments at all. We can, and do, all the time. That is the essence of the task before the judge, for instance, in sentencing criminal offenders or apportioning damages among joint tortfeasors. The point is just that we are *very much* more confident in saying that 'X deserves a longer sentence (or a larger proportion of the damages) than Y' than we are in saying 'X deserves two (or two-and-a-half, or three) times as much as Y'.

¹⁶ True, our inability to calculate relative *ex post* liabilities stands as more of an obstacle to the application of the blame-responsibility model, which after all analyses all responsibilities in terms of such judgments. But the fact that the model is hard to apply in practice does not prove that it is not the morally correct model which we ought try, however imperfectly, to implement in practice.

to the absence of interval-level measures of their relative strengths, thus powerfully suggests that ours is first and foremost a concern with that kind of responsibility that assigns tasks rather than fixes blame.¹⁷

Department of Government,
University of Essex,
Wivenhoe Park,
Colchester CO4 3SQ,
England.

¹⁷ For their comments on earlier drafts of this paper, I am grateful to Diane Gibson, Martin Hollis, Philip Pettit and Jeremy Waldron.